


DSpark audio head — real vLLM serving speedup from the config-loader fix



Qwen3-Omni-30B-A3B Thinker + DSpark audio draft (block 7), LibriSpeech ASR, greedy, 1xB300, 12 utts, identical 669 output tokens. vllm/transformers_utils/configs/speculators/algos.py update_dspark: dspark_bonus_anchor = not config_dict.get("sample_from_anchor", True) (was hardcoded True)